

Political Framing in AI-Generated News Summaries: A Comparative Pilot Study of Five Large Language Models

Akshit Sharma

Age: 17 years

Independent Researcher

Research paper draft prepared with Appendix A: Research Dataset and Coding Workbook

May 2026

Abstract

This study examines whether free and publicly accessible AI chatbots produce meaningfully different summaries when given the same neutral prompt and the same set of politically sensitive news articles. A small pilot dataset was created from six news articles, thirty model-generated summaries, and five widely available AI systems: ChatGPT, Claude, Gemini, Grok, and DeepSeek. Each summary was coded for neutrality, emotional language, political framing, accuracy, and completeness. The results suggest that model behavior varies in ways that matter for public understanding of news. Claude achieved the strongest overall performance in this dataset, while Grok was consistently strong on completeness. ChatGPT often produced balanced prose but sometimes omitted important details, and Gemini and DeepSeek showed more variation depending on source framing. The findings do not prove general model bias because the dataset is intentionally limited. Instead, the study should be read as an exploratory pilot showing that even free public AI tools can reshape news through selection, emphasis, omission, and inherited source framing. The full dataset and coding workbook are provided as Appendix A.

Introduction

AI-generated summaries are becoming a normal part of how people encounter information. A user can paste a news article into a chatbot and receive a clean paragraph in seconds. That convenience matters, especially for younger readers and everyday users who may not have time, paid subscriptions, or the patience to compare multiple sources. At the same time, summarization is not a neutral mechanical act. A summary decides what is important, what can be left out, which words carry the tone, and how much context the reader receives.

This paper studies that problem through a small set of politically sensitive news articles. Although the selected articles focus on immigration-related policy and enforcement events, the broader question is about news in general. Politically charged topics make the issue easier to observe because language choices become visible. A summary can sound professional and still

emphasize danger, victimhood, enforcement, legality, discrimination, or institutional failure. It can also be accurate in one sense while incomplete in another.

The purpose of this study is not to attack or promote any model. The goal is to examine how free and commonly available AI tools behave under the same prompt. The project uses a simple research design: collect news articles from different source orientations, ask multiple AI systems to summarize each article with the same neutral instructions, then code the outputs using a consistent rubric. The resulting workbook is treated as the research dataset for this paper.

Research Questions

- RQ1: Do free or publicly accessible AI models differ in neutrality, accuracy, completeness, emotional language, and political framing when summarizing news?
- RQ2: Does the source orientation of a news article affect how AI summaries frame the same broad topic area?
- RQ3: Which models performed most consistently in this small dataset?

Background

The concern behind this project connects to older research on framing and newer research on large language models. Entman (1993) defines framing as selecting some aspects of reality and making them more salient in communication. That idea is useful for AI summaries because summaries do not simply shorten text. They select, compress, and reorder information.

Recent language-model research has also warned that model outputs can reproduce social patterns, hidden assumptions, and source-language tendencies. Bender et al. (2021) argue that large language models can produce fluent text without genuine understanding, which means readers may mistake polish for reliability. Maynez et al. (2020) show that summarization systems can produce unfaithful or unsupported content, and Weidinger et al. (2021) describe broader social risks from language models, including representational harms and information risks.

This study applies those concerns to public-facing news summarization. The emphasis is practical rather than theoretical: if ordinary users rely on free AI summaries, how much do the summaries differ across tools, and what kinds of differences should users notice?

Methodology

The dataset contains six news articles, thirty AI-generated summaries, and five AI models. The articles were grouped by source orientation into left, center, and right categories for coding purposes. The models were ChatGPT, Claude, Gemini, Grok, and DeepSeek, using free or publicly accessible modes available to a normal user. The study did not use paid enterprise APIs, institutional research infrastructure, or private model access.

Each model received the same neutral summarization instruction. The prompt asked the model to write a 140 to 160 word summary, remain neutral and objective, avoid adding information not in the article, avoid omitting major facts, use professional journalistic language, and return only the summary text. The exact prompt and all outputs are preserved in Appendix A.

Each summary was manually coded across five dimensions. Neutrality was scored from 1 to 5, where a higher score indicates a more neutral summary. Emotional language was scored from 1 to 5, where a higher score indicates more emotionally loaded wording. Political framing was scored from -2 to +2, where negative values indicate left-leaning framing, positive values indicate right-leaning framing, and zero indicates minimal apparent lean. Accuracy and completeness were scored from 1 to 5.

The final quality score used in the workbook is a simple average of neutrality, accuracy, and completeness. Emotional language and political framing were analyzed separately because they measure tone and orientation rather than overall quality. This choice keeps the scoring transparent and easy to audit.

The five coding measures were defined before interpretation: neutrality measured visible opinion or judgment, emotional language measured loaded wording, political framing measured left-to-right emphasis on a -2 to +2 scale, accuracy measured fidelity to the article, and completeness measured coverage of key actors, events, evidence, and context.

Results

The workbook shows that all five models can produce readable summaries, but they do not perform identically. The most important differences appear in completeness, specificity, and the degree to which a model carries forward the framing of the source article. Table 2 summarizes the model-level averages from Appendix A.

Table 1
Average scores by AI model.

Model	Neutrality	Emotional	Framing	Accuracy	Completeness	Avg. Quality
Gemini	3.0	3.0	0.0	3.8	4.2	3.7
DeepSeek	3.2	2.8	0.2	4.3	4.2	3.9
Grok	3.8	2.2	0.3	5.0	5.0	4.6
ChatGPT	3.7	2.0	-0.2	3.7	3.5	3.6
Claude	4.5	1.7	-0.2	4.8	4.8	4.7

Claude had the highest average quality score in the dataset. Its summaries were generally careful about attribution, avoided unnecessary emotional language, and usually preserved the main facts without adding unsupported claims. Grok followed closely, especially because it frequently included more article-specific details than the other models. This made Grok strong on completeness, although it sometimes carried more of the source article's framing into the summary.

ChatGPT often produced smooth and balanced prose, but smoothness did not always equal completeness. In several cases, it omitted names, mechanisms, statistics, or legal details that

other models preserved. Gemini and DeepSeek showed more variation. Some summaries were clear and useful, while others either introduced questionable details, used more loaded language, or leaned more heavily into the article's political frame.

Table 2
Average scores by article.

Article	Neutrality	Emotional	Framing	Accuracy	Completeness	Avg. Quality
L1	3.6	2.4	-0.6	4.2	4.0	3.9
L2	3.2	2.8	-1.2	3.8	4.0	3.7
C1	3.8	2.0	0.2	4.2	4.2	4.1
C2	4.2	1.8	0.0	5.0	4.8	4.7
R1	3.2	2.8	1.2	4.2	4.4	3.9
R2	3.8	2.2	0.6	4.6	4.6	4.3

Article-level results suggest that source framing matters. The center-source articles, especially C2, produced the cleanest cross-model results. The left-source articles tended to produce more negative political framing scores, while the right-source articles tended to produce more positive framing scores. This does not mean the models invented political lean from nothing. Rather, it suggests that models often inherit or amplify the language and emphasis already present in source material.

Table 3
Average scores by source-orientation group.

Source Group	Neutrality	Emotional	Framing	Accuracy	Completeness	Avg. Quality
Left	3.4	2.6	-0.9	4.0	4.0	3.8
Center	4.0	1.9	0.1	4.6	4.5	4.4
Right	3.5	2.5	0.9	4.4	4.5	4.1

Discussion

The central finding is simple but important: AI summaries can sound neutral while still changing how a reader understands the news. The changes are not always dramatic. Often they appear as small choices: which statistic is included, whether a government claim is attributed clearly, whether protesters are described as protesters or rioters, whether a legal process is explained fully, or whether a summary leaves out a counterargument.

This matters because users often judge AI output by fluency. A polished paragraph can create trust even when it is incomplete. In this dataset, the strongest summaries were not merely well-written. They were specific, attributed, and careful about not adding claims beyond the article. The weaker summaries were not always badly written. Some were readable but too generic, while others included details that were not clearly supported by the source article.

The results also show why source framing cannot be ignored. A model summarizing a strongly framed article may reproduce that framing unless it actively neutralizes the language. This is especially visible in politically sensitive news, but the same issue can apply to business, science, crime, health, war, education, or any topic where wording affects public interpretation.

For everyday readers, the practical lesson is that AI summaries should be treated as reading aids, not replacements for the original article. For researchers and students, the lesson is that model comparison should look beyond whether a summary sounds good. Accuracy, completeness, tone, and framing need to be evaluated separately.

Limitations

This project is a pilot study, not a final general theory of AI bias. The dataset is small: six articles and thirty summaries. A larger project would need more articles, more topics, repeated outputs, multiple coders, and stronger statistical testing.

The small size was partly a practical constraint. The study was conducted with free or publicly accessible models, limited hardware, and limited time. Since the researcher is young and working without institutional funding, data collection had to stay realistic. Device heating and limited access to paid tools also affected how much material could be collected. These limits do not make the study useless. They simply mean the findings should be treated as exploratory.

Manual coding is another limitation. Although the rubric was applied consistently, scoring neutrality or political framing always involves judgment. Future versions of the study could use multiple human coders, inter-coder reliability checks, and additional automated text-analysis measures.

Finally, AI systems change over time. A model's behavior in May 2026 may not match its behavior later. The workbook should therefore be read as a snapshot of public-facing model behavior at the time of collection.

Conclusion

This study suggests that free public AI models can summarize news effectively, but they do not all summarize in the same way. Claude and Grok performed strongest overall in this dataset, with Claude showing the best balance of neutrality and precision and Grok showing strong completeness. Other models produced useful summaries but showed more variation in detail, framing, or factual precision.

The broader conclusion is not that AI summaries are bad. The stronger conclusion is that AI summaries are powerful enough to shape understanding. If a reader receives only the summary, then model choices about detail, tone, and framing become part of the news experience. Human review remains necessary, especially when the topic is politically sensitive or when a summary may influence opinion.

References

- Bender, E. M., Gebru, T., McMillan-Major, A., & Shmitchell, S. (2021). On the dangers of stochastic parrots: Can language models be too big? In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency* (pp. 610-623).
<https://doi.org/10.1145/3442188.3445922>
- Entman, R. M. (1993). Framing: Toward clarification of a fractured paradigm. *Journal of Communication*, 43(4), 51-58. <https://doi.org/10.1111/j.1460-2466.1993.tb01304.x>
- Maynez, J., Narayan, S., Bohnet, B., & McDonald, R. (2020). On faithfulness and factuality in abstractive summarization. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics* (pp. 1906-1919). <https://doi.org/10.18653/v1/2020.acl-main.173>
- New York Post. (2026, May 27). Anti-ICE rioters allegedly blasted feds with chemical substance during protest outside Newark's Delaney Hall. <https://nypost.com/2026/05/27/us-news/anti-ice-rioters-allegedly-blasted-feds-with-chemical-substance-during-protest-outside-newarks-delaney-hall/>
- Reuters. (2026a, May 26). Portugal police say operation breaks up large-scale illegal immigration network. <https://www.reuters.com/world/portugal-police-say-operation-breaks-up-large-scale-illegal-immigration-network-2026-05-26/>
- Reuters. (2026b, May 28). Trump administration sues four states for denying ICE undercover license plates. <https://www.reuters.com/legal/government/trump-administration-sues-four-states-denying-ice-undercover-license-plates-2026-05-28/>
- The Guardian. (2026, May 26). US immigration enforcement flights generate surge in carbon emissions. <https://www.theguardian.com/us-news/2026/may/26/us-immigration-flights-emissions>
- The Times of India. (2026, May 30). Canada's new immigration policy could let foreign criminals delay deportation: What we know.
<https://timesofindia.indiatimes.com/world/rest-of-world/canadas-new-immigration->

policy-could-let-foreign-criminals-delay-deportation-what-we-know/articleshow/131284030.cms

The Verge. (2026, May 27). Trump's ICE relies on racial profiling for mass deportations, report finds. <https://www.theverge.com/policy/939434/trump-ice-racial-profiling-dhs-mass-deportations>

Weidinger, L., Mellor, J., Rauh, M., Griffin, C., Uesato, J., Huang, P. S., Cheng, M., Glaese, A., Balle, B., Kasirzadeh, A., Biles, C., Brown, S., Kenton, Z., Hawkins, W., Stepleton, T., Birhane, A., Haas, J., Rimell, L., Hendricks, L. A., ... Gabriel, I. (2021). Ethical and social risks of harm from language models. arXiv. <https://arxiv.org/abs/2112.04359>

Appendix A: Research Dataset and Coding Workbook

The full dataset, article metadata, model summaries, scoring sheet, and dashboard are provided in the accompanying workbook:

<https://github.com/FreakyDevelopers/Political-Framing-in-AI-News-Summaries/blob/main/data/research-dataset.xlsx>

This workbook is part of the research record. It contains the article list, the exact neutral prompt used for model outputs, all thirty AI-generated summaries, manual coding notes, and the summary dashboard used for the results tables in this paper.